

Using Iris for Program Verification

Abel Nieto

March 27, 2020

1 Monotone Counter

▷ **Coq code:** <https://github.com/abeln/iris-practice/blob/master/counter.v> ◁

This is the counter example from Section 7.7 of the notes. The module exports three methods:

- **newCounter** creates a new counter. Counters are represented simply with an int reference.

```
Definition newCounter : val := \lambda: <>, ref #0.
```

- **read** returns the value currently stored in a counter.

```
Definition read : val := \lambda: "c", !"c".
```

- **incr** takes a counter, increments it, and returns unit.

```
Definition incr : val :=  
  rec: "incr" "c" :=  
    let: "n" := !"c" in  
    let: "m" := #1 + "n" in  
    if: CAS "c" "n" "m" then #() else "incr" "c".
```

The client for the counter is a program that instantiates a counter, spawns two threads each incrementing the counter, and then reads the value off the counter:

```
Definition client : val :=  
  \lambda: <>,  
    let: "c" := newCounter #() in  
    ((incr "c") ||| (incr "c")) ;;  
    read "c".
```

1.1 Specs

We can prove two different specs for the code above:

- Authoritative RA

In the first spec, our resource algebra is $\text{AUTH}(\mathbb{N})$. Elements of this RA are of the form either $\bullet n$ (authoritative) or $\circ n$ (non-authoritative), where $n \in \mathbb{N}$.

The important properties of this RA (for the counter example) are:

- $\bullet 0 \cdot \circ 0 \in \mathcal{V}$
- $\bullet m \cdot \circ n \in \mathcal{V}$ implies $m \geq n$
- $\bullet m \cdot \circ n \rightsquigarrow \bullet(m+1) \cdot \circ(n+1)$
- $|\circ n| = \circ n$

Then we can define our predicate for the counter:

$$\text{isCounter}(l, n, \gamma) = \ulcorner \circ \bar{n} \urcorner^\gamma * \exists s. \boxed{\exists m. l \hookrightarrow m \wedge \ulcorner \bullet \bar{m} \urcorner^\gamma}^s$$

This predicate is persistent because the invariant component is persistent and $\circ n$ is also persistent (via the **PERSISTENTLY-CORE** rule), because the core of $\circ n$ is $\circ n$ itself.

The intuition behind the predicate is that if we own $\text{isCounter}(l, n, \gamma)$ then, at an atomic step (e.g. a load), we can know that l is allocated and it points to a value that's $\geq n$.

Additionally, after an increment, we can update the resource to $\text{isCounter}(l, n+1, \gamma)$.

We then prove the following specs for the counter methods:

- $\{\top\} \text{newCounter } () \{l. \exists \gamma. \text{isCounter}(l, 0, \gamma)\}$
- $\{\text{isCounter}(l, n, \gamma)\} \text{read } l \{v. v \geq n\}$. Notice how we don't have to give the counter back, because it's persistent.
- $\{\text{isCounter}(l, n, \gamma)\} \text{incr } l \{\text{isCounter}(l, n+1, \gamma)\}$

We can also prove the client spec:

$$\{\top\} \text{client } () \{n. n \geq 1\}$$

The “problem” with this specification *isCounter* doesn't tell us anything about what other threads are doing with the counter. So in the verification of the client code, after both threads return, we know

$$\text{isCounter}(l, 1, \gamma) * \text{isCounter}(l, 1, \gamma) \equiv \text{isCounter}(l, 1, \gamma)$$

So we know that the counter must be ≥ 1 , but we don't know that it's 2.

- Authoritative RA + Fractions

The second spec we can give is more precise, and allows us to conclude that the client reads exactly 2 once the two threads finish.

The resource algebra we use is $\text{AUTH}((\mathbb{Q}_p \times \mathbb{N})_?)$:

- \mathbb{Q}_p is the RA of positive fractions, where the valid elements are those ≤ 1 .
- The $?$ is the *optional RA* construction. In this case, it's needed because to use the authoritative RA we need the argument algebra to be *unital*. And $\mathbb{Q}_p \times \mathbb{N}$ is not unital because 0 is not an element of \mathbb{Q}_p .

The important properties of this RA are:

- $\circ(1, 0) \cdot \bullet(1, 0) \in \mathcal{V}$
- $\circ(p, n) \cdot \bullet(1, m) \in \mathcal{V}$ implies $m \geq n$
- $\circ(1, n) \cdot \bullet(1, m) \in \mathcal{V}$ implies $m = n$. This is the rule that will allow us to give a more precise specification.
- $\circ(p, n) \cdot \bullet(q, m) \rightsquigarrow \circ(p, n+1) \cdot \bullet(q, m+1)$

With this RA, we get the new counter resource

$$\text{isCounter}(l, n, p, \gamma) = \left[\begin{array}{c} \text{---} \\ \circ(p, n) \\ \text{---} \end{array} \right]^\gamma * \exists s. \left[\begin{array}{c} \exists m. l \hookrightarrow m \wedge \left[\begin{array}{c} \text{---} \\ \bullet(1, m) \\ \text{---} \end{array} \right]^\gamma \end{array} \right]^s$$

This resource is no longer persistent, because the core $|p|$ is undefined. However, we can show the following

$$\text{isCounter}(l, n+m, p+q, \gamma) \dashv\vdash \text{isCounter}(l, n, p, \gamma) * \text{isCounter}(l, m, q, \gamma)$$

We then prove the following specs for the counter methods:

- $\{\top\} \text{newCounter } () \{l. \exists \gamma. \text{isCounter}(l, 0, 1, \gamma)\}$
- $\{\text{isCounter}(l, n, p, \gamma)\} \text{read } l \{v. v \geq n * \text{isCounter}(l, n, p, \gamma)\}$. Notice how we need to give back the counter in this case.
- $\{\text{isCounter}(l, n, 1, \gamma)\} \text{read } l \{v. v = n * \text{isCounter}(l, n, p, \gamma)\}$. Notice that since we have the entire fraction, we can get the exact value of the counter.
- $\{\text{isCounter}(l, n, p, \gamma)\} \text{incr } l \{\text{isCounter}(l, n+1, p, \gamma)\}$

With this, we get the more precise client spec

$$\{\top\} \text{client } () \{n. n = 2\}$$

2 Locks and Coarse-Grained Bags

▷ Coq code: <https://github.com/abeln/iris-practice/blob/master/lock.v> ◁

This is the spin lock example from Section 7.6 of the notes.

The spin lock is a module with three methods:

- **newlock** creates a new lock. Locks are represented as a reference to a boolean. If **false**, the lock is free; if **true**, then it's taken.
- **acquire** uses a CAS cycle to spin until it can acquire the lock.
- **release** just sets the lock to **false**.

The code is as above:

```
Definition newlock : val := \lam: <>, ref #false.
```

```
Definition acquire : val :=
  rec: "acquire" "l" :=
    if: CAS "l" #false #true then #() else "acquire" "l".
```

```
Definition release : val := \lam: "l", "l" <- #false.
```

2.1 Spec

The involved RA is $\text{EX}(\text{UNIT})$.

- Notice that unlike other RA constructions, $\text{EX}(S)$ just requires S to be a *set*, as opposed to another RA.
- The exclusive RA is defined by adding an element \perp to the carrier set. For any two elements $x, y \neq \perp$, we have $x \cdot y = \perp$. Every element except for \perp is valid.
- What this means is that if we own $\llbracket \tilde{x} \rrbracket^\gamma$, no other thread can own another $\llbracket \tilde{y} \rrbracket^\gamma$, because their product would be invalid.

The lock predicate is then

$$\text{isLock}(l, P, \gamma) = \exists s \left[l \hookrightarrow \text{false} \wedge P \wedge () \vee l \hookrightarrow \text{true} \right]^s$$

A first observation is that the lock protects an *arbitrary predicate* P , which is what the client uses to prove its correctness.

The intuition is the following:

- If the lock is *unlocked* ($l \hookrightarrow \text{false}$), then we need to *give away* to the invariant the predicate P protected by the lock *and* the key $()$. Because $() \in \text{EX}(\text{UNIT})$, we *know that no other key can be created*.
- When the lock is *locked* ($l \hookrightarrow \text{true}$), we don't need to give away any resources.

- When the lock transitions from unlocked to locked, we *gain* resources.
- When the lock transitions from locked to unlocked, we *lose* resources.

The last two points are reflected in the specs for the lock methods:

- $\{P\}\text{newlock } ()\{l.\exists\gamma.\text{isLock}(l, P, \gamma)\}.$

Notice that even though only one element of $\text{EX}(\text{UNIT})$ is allowed, this restriction is *per location/ghost name*, so we're allowed to allocate a new (exclusive) ghost resource, provided we furnish a new ghost name γ .

- $\{\text{isLock}(l, P, \gamma)\}\text{acquire } l\{P * ()\}$

After acquiring the lock, we gain ownership of the predicate P and the key $()$.

- $\{\text{isLock}(l, P, \gamma) * P * ()\}\text{release } l\{\top\}$

To release the lock, we need to show that P continues to hold, and that we have the key $()$.

Also note that the lock predicate is *persistent*, since it's protected inside an invariant. This is important because multiple threads need to know that a given lock exists (so they can synchronize).

2.2 Client: Coarsed-Grained Bags

This client is also from the notes, and it consists of a *bag* data structure. We can create new bags, add items to it, and remove (the first) items from it.

▷ **The bag looks just like a queue** ◁.

The bag code follows:

Section `bag_code`.

```
(* A bag is a pair of (optional list of values, lock) *)
Definition new_bag : val :=
  \lam: <>, (newlock #(), ref NONEV).
```

```
(* Insert a value into the bag. Return unit. *)
Definition insert : val :=
  \lam: "bag" "v",
    let: "lst" := Snd "bag" in
    let: "lock" := Fst "bag" in
    acquire "lock";;
    "lst" <- SOME ("v", !"lst");;
    release "lock";;
    #().
```

```
(* Remove the value last-added to the bag (wrapped in an option), if the
```

```

    bag is non-empty. If the bag is empty, return NONE. *)
Definition remove : val :=
  \lam: "bag",
    let: "lst" := Snd "bag" in
    let: "lock" := Fst "bag" in
    acquire "lock";;
    let: "res" :=
      match: !"lst" with
      NONE => NONEV
      | SOME "pair" =>
        "lst" <- Snd "pair";;
        SOME (Fst "pair")
      end
    in
    release "lock";;
    "res".

```

End bag_code.

2.2.1 Bag Specification

The spec we want to show says only elements satisfying some predicate Φ can be added to the bag. In turn, we guarantee that any element x removed from the bag satisfies Φx .

The bag predicate looks like

$$\text{isBag}(b, \Phi, \gamma) = \exists l, lst. b = (lock, lst) \wedge \text{isLock}(l, \exists n, vs. lst \hookrightarrow vs \wedge \text{baglist}(vs, n, ,)\gamma)$$

In turn, $\text{baglist}(vs, n, \Phi)$ is a predicate that expresses that vs is a `HeapLang` list of n elements, all of which satisfy Φ .

$$\begin{aligned} \text{baglist}(vs, n, \Phi) = & (n = 0 \wedge vs = \text{None}) \\ & \vee (n = Sn' \wedge \\ & \exists v, vs'. vs = \text{Some}((v, vs')) \wedge \Phi v \wedge \text{baglist}(vs, n', \Phi)) \end{aligned}$$

▷ The notes use guarded recursion to define `baglist`. I ended using the index n . Need to try the guarded recursion approach. ◁

The specs for the different methods are:

- $\{T\} \text{newbag } () \{v. \gamma \text{isBag}(v, \Phi, \gamma)\}$
- $\{\Phi v * \text{isBag}(b, \Phi, \gamma)\} \text{insert } b \ v \{T\}$
- $\{\text{isBag}(b, \Phi, \gamma)\} \text{remove } b \{v. v = \text{None} \vee \exists x. v = \text{Some}(x) * \Phi x\}$

Notice that `isBag` is also persistent/duplicable.

▷ Need to implement concurrent client for bags ◁